

peer-to-peer and agent-based computing

Programming P2P Applications




Plan of lecture

- Parallel programming
- Demonstrators and prototypes
- Message-passing vs. memory-sharing
- Tuple spaces as an alternative
- Threads with message-passing



Parallel Programming

- Traditional computing is **serial**
- However, some problems can be naturally broken down into smaller problems which can be tackled in parallel
- Example: processing a pay-roll



Parallel Programming (2)

- Multi-processor devices are very common
 - 2-core, quad-core, eight-core, ...
 - Applications do not need to “take turns” when run
 - Real parallelism!
 - Parallel computing is now very important
- Different kinds of parallel programs:
 - Processes **share variables** (concurrent programming)
 - Process **send messages** (distributed programming)To share info/data, coordinate, etc.

Why should we care about this?

- P2P applications are different from conventional computing
 - Message-passing is the name of the game
 - Not just HTTP requests (one-off messages)
 - Starting requests, forwarding messages, etc.
- However, if we want to “play” with P2P apps, then we might need to **simulate** things
 - For instance, start up 5,000 peers to see how it goes...
 - Message-passing can be an issue (more on this later)
- Let’s look at features of P2P applications and options to help us implement these

Important features of P2P applications

From an **individual peer**’s perspective:

- Loosely-coupled message-passing
 - No “blocking” command to read message
- Messages may not arrive in the order you want or need

From the **global (system)** perspective:

- High number of peers (50 and beyond)
- Very high volume of messages

Questions:

- How many machines can we use?
- Is our prototype to be deployed later?
- Is it for simulation only (to be re-developed later)?

Demonstrators and prototypes

- Suppose you have a very good idea for a P2P app
 - A protocol (order and kinds of messages)
 - An information model (routing tables, info in messages)
 - Some decision process for peers to follow
- What should you do?
 - We need to get to the core aspects quickly, to check if our ideas are any good
 - A demonstrator/prototype will serve this purpose
 - No GUIs, no sound effects, just core functionalities
 - Once we are satisfied that our ideas work, then an actual implementation could be developed
 - We might need to throw away initial prototype...
 - ... Hence it cannot be too sophisticated!

Demonstrators and prototypes (2)

- The main choices of technologies for a prototype:
 - Actual message-passing (distributed computing)
 - Simulation of message-passing (concurrent computing)
- Actual message-passing:
 - **Advantages:**
 - Realism (taste of what you will use for the final app)
 - Many machines can be used when testing
 - **Disadvantages:**
 - If you have only one machine, message-passing can be costly
 - Low-level protocols (HTTP, TCP) may delay, lose or corrupt messages
- Simulation of message-passing
 - Pretty much the reverse situation
 - NB: ideal when only one machine to be used

Message-passing

- Basic socket and ports
 - Programming language must offer this feature
- Decision: messages passed around as
 - Strings (which need to be parsed to be processed) OR
 - Serialised objects (no need to parse)
- Non-blocking receive
 - Programming language must offer this feature
- Single vs. many-threaded:
 - Each peer is a stand-alone program
 - All peers are threads in a single program
 - Peers are threads, but many copies of a program running
 - Peers themselves can be multi-threaded

Message-passing (2)

- Example: P2PTool
 - In-house educational tool (not for actual deployment)
 - Scale-up (more than 20 peers) is error-prone
 - Messages are serialised objects
 - Peers running in many machines can join network
- Further complication:
 - To provide a “global” view of the network, all messages are “intercepted” by the server
 - This puts a lot of strain on one process
 - Messages go missing and get corrupted

Simulating message-passing


- Peers are threads running in one same program
 - Peers themselves could be multi-threaded...
 - Threads sharing working memory (RAM)
- Messages are managed via **a shared data structure**
 - An array, where all messages are stored
 - All peers/threads have access to shared data structures
- Challenge: discipline peers’ access to messages
 - Being a shared variable, we must ensure **fairness**
 - For instance, each peer operates for n ms (round-robin)

Simulating message-passing (2)

- PeerSim (<http://peersim.sourceforge.net/>) simulates message-passing
 - Shared data structure
 - Multi-threaded (each peer is a thread?)
 - Peers take turns accessing this data structure
- Only one JVM (Java Virtual Machine)
 - Very efficient (complex scenarios quickly modelled)
 - Scalable (10,000 peers and beyond)
 - Various levels of control (useful in simulations)

An alternative: tuple spaces

- Tuple spaces
 - Processes write onto and read from shared tuple space
 - Tuple: pair, triple, quadruple,..., of pieces of information
 - A kind of “data-base” (centralised) where info is kept
- Associative memory:
 - Use of patterns to access tuples – anything which matches the pattern is returned




An alternative: tuple spaces (2)

- Programming languages offer extensions
 - JavaSpaces for Java
- Infrastructure to start up & manage tuple spaces
- Extensions to
 - Write onto tuple space
 - Read from tuple space
 - Remove tuples from space
- Remotely, with exceptions
- Interesting functionality:
 - Express interest on a tuple pattern and be warned when it arrives
 - Better than that: a method is invoked in your program, when a tuple arrives

An alternative: tuple spaces (3)

- Programming with tuple spaces is fun...
- Example: ensure 3 processes all read tuple
 - Assume you “own” all processes
 - Design info in tuple to help the solution



JADE

- Java Agent Development Environment
 - <http://jade.tilab.com/>
- Agents \approx peers
 - Agents are lightweight Java threads
 - Messages sent/received within one JVM
- However,
 - It is possible to have JADE copies running in various devices (mobile phones!)
 - Many agents in each of them
 - Agent in a device can send message to agents in other devices
 - Agent can migrate (move) from one device to another
 - Mobile agents
